

Rust programming language

a systems language

pursuing the trifecta

safe, concurrent (?), fast

Introduction to safe systems programming

Quick introduction

- Mozilla Research project
- Started by Graydon Hoare in 2006
- 1.0-alpha just released, 1.0 set to be released in **early March**
- Main niches:
 - Large scale maintainable systems: browser engines, games, databases
 - Fast, *predictably fast* code
 - Correctness and confidence in the quality
- Substantial influence of functional languages
- Still, striving to be *close to the machine*

```
fn main() {  
    println!("Hello world!");  
}
```

Run

```
fn main() {
    // A simple integer calculator:
    // + or - means add or subtract by 1
    // * or / means multiply or divide by 2

    let program = "+ + * - /";
    let mut accumulator = 0i;

    for token in program.chars() {
        match token {
            '+' => accumulator += 1,
            '-' => accumulator -= 1,
            '*' => accumulator *= 2,
            '/' => accumulator /= 2,
            _ => { /* ignore everything else */ }
        }
    }

    println!("The program \"{}\" calculates the value {}", program, accumulator);
}
```

Run

What choice do programmers have today?

Close to the metal but easy to get wrong High level, secure but with less control

- C
- C++

- Java
- Scala
- Python
- Ruby
- JavaScript
- Haskell



Control

 Emmanuel Huybrechts

- Deterministic resource management: *memory, sockets, locks*
 - Resource Acquisition Is Initialisation

```
use std::io::File;
fn main() {
    let file = File::open(&Path::new("/Users/jakub/TODO.txt"));
    // file closed when it goes out of scope.
}
```

Run

- Machine-sympathetic memory layout

```
use std::intrinsics::transmute;
struct Point { x: u32, y: u32 }
struct Rectangle {
    top_left: Point,
    width: u32,
    height: u32
}
fn main() {
    let rectangle = Rectangle {
        top_left: Point { x: 0, y: 0 },
        width: 50,
        height: 50
    };

    println!("{}\n{}", transmute(&rectangle.top_left.y as *const u32 as u64),
        transmute(&rectangle.width as *const u32 as u64));
}
```

Run

- Zero-overhead memory abstractions

```
struct Point {  
    x: u32,  
    y: u32  
}  
struct Rectangle {  
    top_left: Point,  
    width: u32,  
    height: u32  
}  
fn main() {  
    println!("{}", std::mem::size_of::<Rectangle>());  
}
```

Run

- Can use on bare metal (no runtime).
- Can interoperate with C libraries.
- Statically determined memory deallocation.

```
fn print_and_forget(x: String) {  
    println!("{}", x);  
    // free(x);  
}
```

Run

Safety



Classes of errors common in C++

- Dangling pointers
- Out-of-bounds array accesses
- Memory leaks
- Buffer overrun
- Use-after-free
- Data races
- Iterator invalidation

Classes of errors common in Java

- NullPointerException
- ConcurrentModificationException
- Out of heap



Security

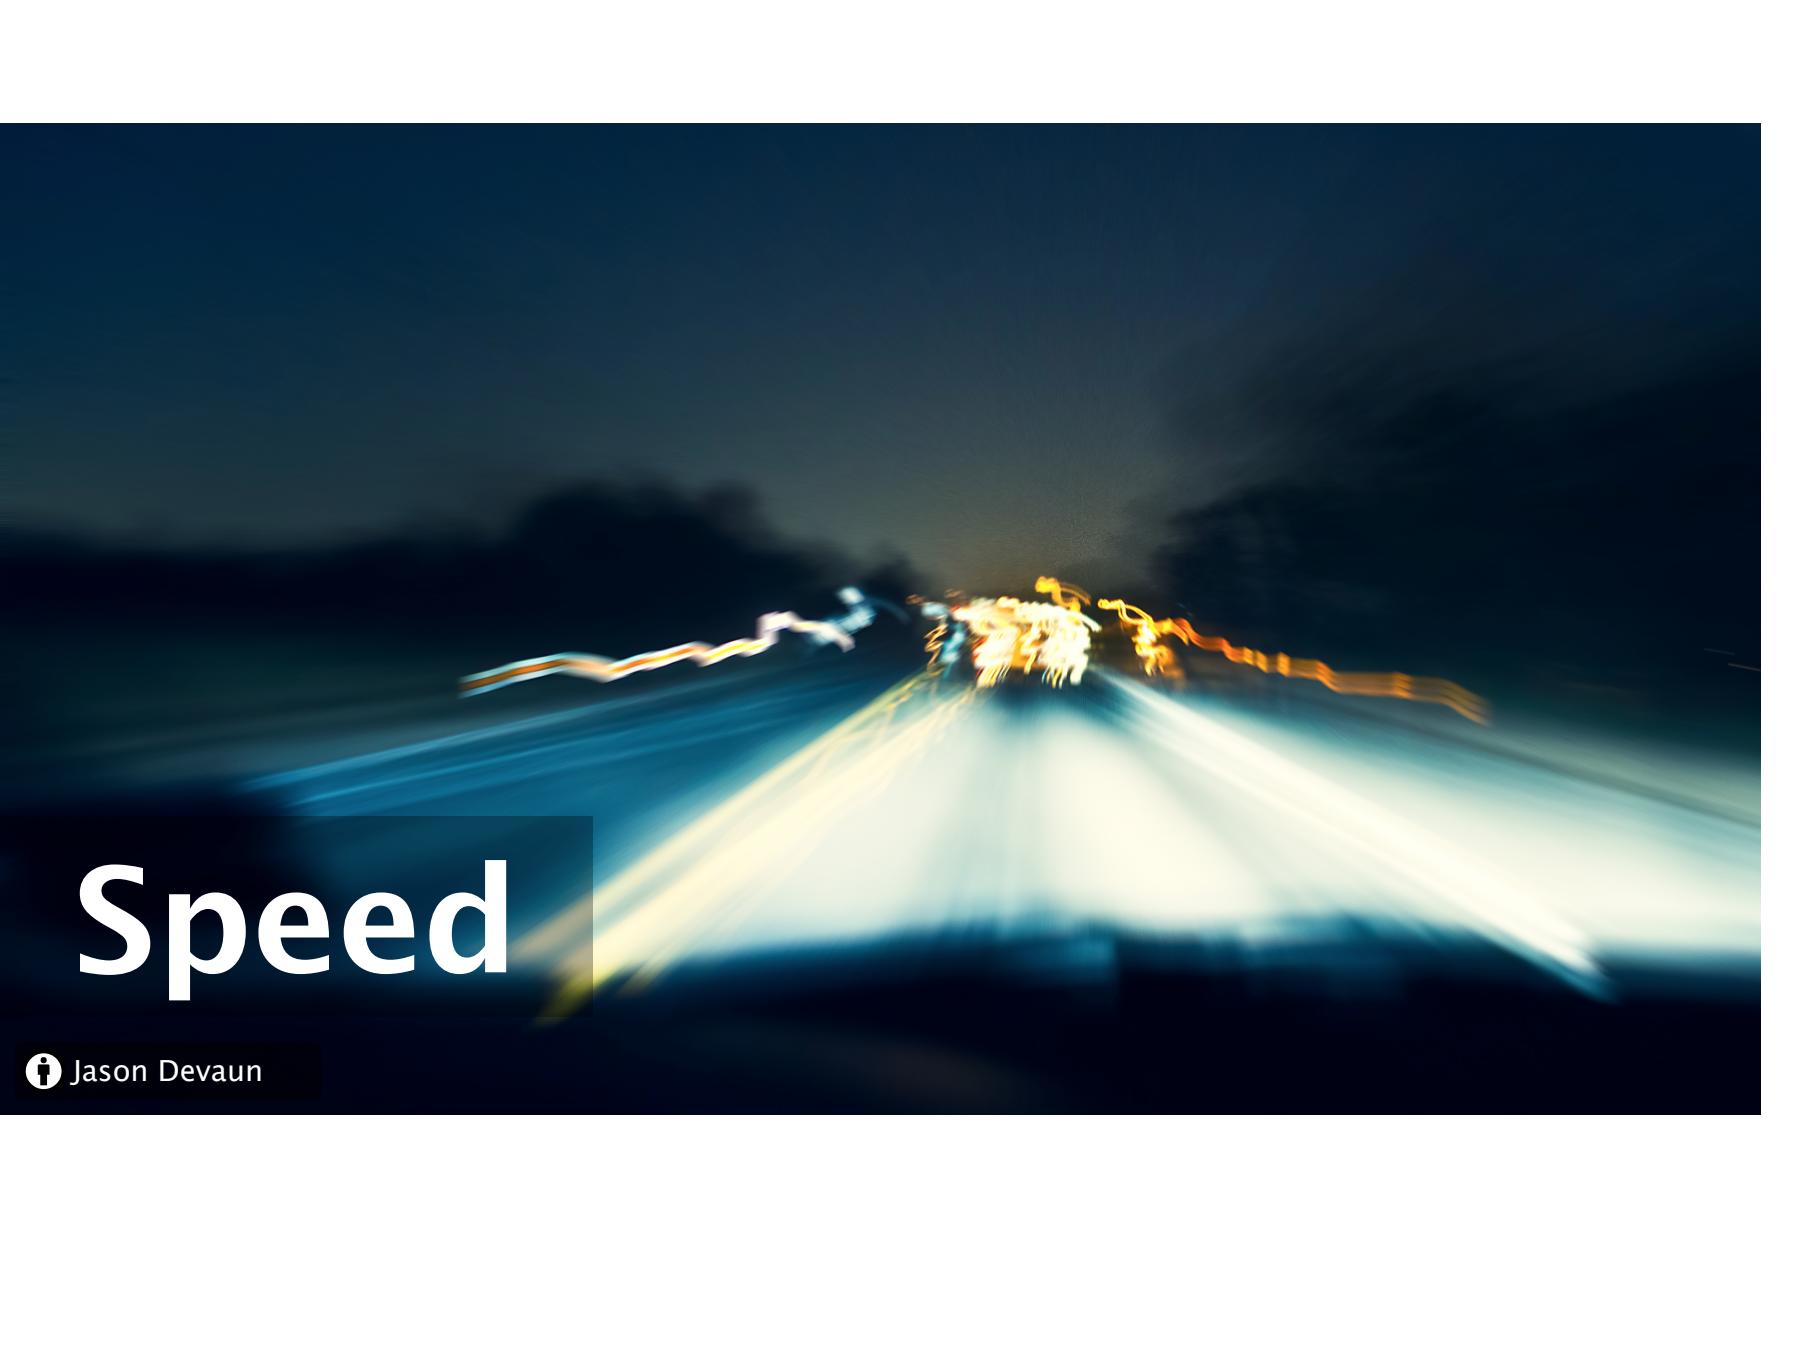
 Rob

Security hazards

Security Advisories For Firefox

We don't want any of this.

- *Best practices* are not enough.
- Enforcing safety via idioms is not working.
- The compiler should reject programs that are not safe.



Speed

 Jason Devaun

Performance

- Very good performance – on-par with C and C++
 - State-of-the-art optimization pipeline using LLVM
- *Predictable* performance
 - Ahead-of-time compilation to a static binary
 - No garbage collection
 - Static dispatch whenever possible
- Very early days
 - LLVM driven with C, C++ frontends in mind
 - The Rust compiler knows more about your programs
- [Computer Language Benchmarks Game](#)

A close-up, shallow depth-of-field photograph of a chocolate cake. The cake is covered in a thick layer of dark chocolate frosting that is dripping down the sides. It is topped with numerous small, colorful, round candies in shades of red, yellow, green, blue, and white. The background is blurred, showing more of the same cake in the distance.

Can we have the cake and eat it too?

 Jason Devaun

Yes.

- Rust guarantees memory safety with strong compile-time analysis.
- Strong static type system
- No explicit memory management, the compiler knows *precisely* when memory can and must be deallocated.
- If a Rust program compiles, it's memory safe.
- Zero-cost abstraction characteristics of a C++ program.
- Type inference makes programs easy to write and read.

Main idea

Ownership

- Who is allowed to release a resource? *Memory*, in particular.
- Which thread can safely access a resource?
- When can a resource be safely *moved* in memory?

```
fn do_things(x: Box<uint>) {  
    println!("{} + {} = {}", x, *x + *x);  
}  
  
fn main() {  
    let foo = box 42u;  
    do_things(foo);  
    println!("{}", foo);  
}
```

Run

- A datum is owned by one owner at a time.

Borrowing

- Who has read access to a resource but doesn't own it?
- Who has mutable access to a resource?
- Cannot lend a mutable reference and an immutable reference at the same time.

```
fn main() {  
    let mut v = vec![0u, 1, 2, 3];  
    let last = &v[3];  
    v.clear();  
    println!("{}", last); // Dangling pointer?  
}
```

Run

- References cannot outlive the scope of a value.

```
fn main() {  
    println!("Your number was: {}", get_number());  
}  
  
fn get_number<'a>() -> &'a int {  
    let a = 3;  
    return &a;  
}
```

Run

Mutability

- When can a resource be modified?

```
fn main() {  
    let x = 42u;  
    x = 43u;  
}
```

Run

```
struct Point { x: uint, y: uint }  
  
fn f(p: &Point) {  
    p.y = 42;  
}  
fn main() {  
    let p = Point { x: 0, y: 0 };  
    f(&p);  
}
```

Run

The best thing

- Not just annotations for the compiler.
- They **describe** the system.
- They **document** the system.

```
struct WebRequest {  
    connection: postgres::Connection,  
    store: &SessionStore  
}
```

Run

- They **help me** reason about the system.
- The compiler catches bugs in the code.
- ...but also in my thinking.

Other features

Pattern matching

- Allows deep inspection of values.
- Allows behaviour conditional on the shape of data.

```
#[deriving(Show)]
enum Direction {
    North,
    East,
    South,
    West
}

fn reverse(dir: Direction) -> Direction {
    match dir {
        North => South,
        East  => West,
        South => North,
        West   => East,
    }
}

fn main() {
    println!("You should head: {}", reverse(North));
}
```

Run

Traits

```
trait HasSize {
    fn size(&self) -> uint;
}
struct Square {
    side: uint
}
impl HasSize for Square {
    fn size(&self) -> uint {
        self.side * self.side
    }
}
struct Rectangle {
    width: uint,
    height: uint
}
impl HasSize for Rectangle {
    fn size(&self) -> uint {
        self.width * self.height
    }
}
fn print_size<T: HasSize>(v: T) {
    println!("The shape has size: {}", v.size());
}
fn main() {
```

Run

#[derive]

- Automatically derive useful traits.

```
struct Point {  
    x: uint,  
    y: uint  
}  
  
fn is_zero(p: &Point) -> bool {  
    *p == Point { x: 0, y: 0 }  
}  
fn main() {  
    assert!(!is_zero(&Point {  
        x: 2, y: 4  
    }));  
}
```

Run

Macros

- Meta-programming with macros "by example"
- Allows introducing convenient simplifications of programs

```
#![feature(macro_rules)]  
  
macro_rules! assert_eq(  
    ($given:expr, $expected:expr) => ({  
        match (&($given), &($expected)) {  
            (given_val, expected_val) => {  
                if (*given_val != *expected_val) {  
                    panic!("{} != {}", *given_val, *expected_val);  
                }  
            }  
        }  
    })  
;  
  
fn main() {  
    assert_eq!(4u, 3u);  
}
```

Run

Parallelism

- Designed with *safe* parallelism in mind
 - Ownership semantics guarantee absence of data races
- Built-in basic support for SIMD

```
use std::iter::MultiplicativeIterator;
use std::sync::Future;
use std::slice::SliceConcatExt;

fn factorial(x: uint) -> uint {
    std::iter::range_inclusive(1, x).product()
}

fn main() {
    println!("{}",
        range(0, 10)
            .map(|n| Future::spawn(move || factorial(n)))
            .map(|mut f| f.get().to_string())
            .collect::<Vec<_>>()
            .len());
}
```

Run

Concurrency

- Message passing-based communication via channels

```
fn main() {
    let (tx, rx) = channel();
    spawn(proc () {
        tx.send(2u + 2);
        tx.send(3u + 5);
    });
    println!("{} , {}", rx.recv(), rx.recv());
}
```

Run

- ...but many other concurrency primitives exposed with safe interfaces

Iterators

- Lazy.
- A major abstraction in libraries.

```
struct Town {  
    population: uint,  
    name: &'static str  
}  
fn show_most_populated_town_under_450k(towns: Vec<Town>) {  
    println!("Most populated town under 450k is {}", towns.iter()  
        .filter(|town| town.population < 450 * 1000)  
        .max_by(|town| town.population)  
        .map(|town| town.name));  
}  
fn main() {  
    let towns = vec![  
        Town { name: "Gothenburg", population: 491630 },  
        Town { name: "Malmö", population: 278523 },  
        Town { name: "Stockholm", population: 789024 },  
        Town { name: "Kärklskrona", population: 35212 },  
        Town { name: "Lund", population: 82800 }  
    ];  
    show_most_populated_town_under_450k(towns);  
}
```

Run

- Compile down to very efficient code, usually equivalent to hand-rolled loops.

Built-in testing and benchmark harness

```
use std::iter::MultiplicativeIterator;
fn factorial(n: uint) {
    range(1, n + 1).product()
}

#[test]
fn test_factorial() {
    assert_eq!(factorial(1), 1);
    assert_eq!(factorial(4), 24);
}

const BENCH_SIZE: uint = 20;
#[bench]
fn iterative_fibonacci(b: &mut Bencher) {
    b.iter(|| {
        fibonacci_sequence().take(BENCH_SIZE).collect::<Vec<uint>>(
    })
}
```

Run

And if you really need it...

Unsafe

```
fn add_two_numbers_and_print(x: *const u64, y: *const u64) {
    println!("{} + {} = {}", *x, *y, *x + *y);
}

fn main() {
    let x = 1 as *const u64;
    let y = 2 as *const u64;
    add_two_numbers_and_print(x, y);
}
```

Run

And if you really need it...

Unsafe

```
fn add_two_numbers_and_print(x: *const u64, y: *const u64) {
    println!("{} + {} = {}", *x, *y, *x + *y);
}

fn main() {
    let x = 1 as *const u64;
    let y = 2 as *const u64;
    add_two_numbers_and_print(x, y);
}
```

Run

But you probably won't.

```
$ cd cargo
$ cat **/*.rs | wc -l
24700
$ git grep unsafe src | wc -l
0
```

```
$ cd rubigo-luculenta
$ cat **/*.rs | wc -l
2999
```

Servo

Research project to build the next-generation Web browser

Goals: security, mobile performance and power efficiency



Cargo package manager

- Reproducible builds.
- Central registry of crates and cargo publish coming soon.

```
[package]
name = "postgres"
version = "0.0.0"
authors = [ "Steven Fackler <sfackler@gmail.com>" ]

[lib]
name = "postgres"
path = "src/lib.rs"
test = false

[dependencies.openssl]
git = "https://github.com/sfackler/rust-openssl"

[dependencies.phf]
git = "https://github.com/sfackler/rust-phf"

[dependencies.phf_mac]
git = "https://github.com/sfackler/rust-phf"
```

Projects (known to be) using Rust

- Servo browser engine
- The Rust compiler
- Cargo package manager
- Piston game engine
- Zinc close-to-metal stack
- Iron Web framework
- Teepee
- Tilde
- OpenDNS

1.0 – what's being fleshed out

- Standard library stability
- ~~Dismantling of M:N concurrency~~
- ~~Moving out many built-in crates out to Cargo packages~~
- ~~Fleshing out terminology for spreading the good news~~
- Language features blocking the release:
 - Associated types
 - Unboxed closures
 - ~~Polished macro system~~
- Crate registry for Cargo (<http://rubygems.org> anyone?)

Rust 2 years ago and now

```
enum Direction {
    North,
    East,
    South,
    West
}

fn main() {
    let mut table = ~[];
    for [0, ..10].each |i| {
        let mut row = ~[];
        for [0, ..128].each |j| {
            row.push(North);
        }
        table.push(copy row);
    }
    io::print(fmt!("before: %d\n", table[0][0] as int));
    setT(&mut table);
    io::print(fmt!("after: %d\n", table[0][0] as int));
}

fn setT(table: &mut ~[~[Direction]]) {
    for [0, ..10].each |i| {
        for [0, ..128].each |j| {
```

Run

1.0 and stability

- Rust known for its backwards-incompatible changes in the 0.x cycle
- 1.0 will come with a *stability* guarantee
 - The core parts of the language will no longer change
 - Some parts of the standard library will no longer break
 - Feature gates will be used for non-stable parts of the language
 - A stable channel will only contain the stable parts of the language
 - A nightly channel will ship the unstable parts of the language
- The language will continue to develop
 - Rapid 6-week release cycle
 - <http://blog.rust-lang.org/2014/09/15/Rust-1.0.html>

Getting started

- Amazing community. User groups in:
 - California, London, Paris, Zurich, Malmö, Seattle, Toronto, Lisbon, Berlin, Tokyo, Milano and many more.
- 1230 Rust repositories on GitHub
- [The Rust Book](#)
- IRC: irc.mozilla.org [#rust](#)
- [This Week in Rust](#)
- <http://github.com/rust-lang>
- <http://doc.rust-lang.org>
- Playpen: <http://play.rust-lang.org>
- Rust By Example: <http://rustbyexample.com>
- Writing a path tracer in Rust: <http://ruudvanasseldonk.com/>

Questions?

jakub@jakub.cc
@JAW_dropping